

# How to use HLS optimization directives for Image Histogram Equalization

**Daniele Bagni,  
DSP Specialist for EMEA**

# Outline

- **The most frequently adopted directives**
- **Some directives suitable for arrays**

# The most frequently adopted directives

# The application: image histogram equalization

- MATLAB: `y_eq = histeq(inp_y, 255);`

input Y



Y equalized



- See Freescale' application note: [AN4318.pdf](#), rev. 0, Jun 2011

# We will see...

## ▪ Step after step how the various directives affect performance

1. `set_directive_loop_tripcount -min 640 -max 1920 func_name/Lx`  
`set_directive_loop_tripcount -min 480 -max 1080 func_name/Ly`
  - To help the compiler optimize an unbounded loop
2. `set_directive_dataflow top_func_name`
  - To make the 3 subroutines concurrent (improves latency and throughput)
3. `config_dataflow -default_channel fifo -fifo_depth 2`
  - To change the default interface of dataflow from ping-pong RAM to FIFO (all in all video data are streamed) buffering
4. `set_directive_pipeline`
  - To pipeline the various routines (improves latency and throughput)
5. `set_directive_dependence false`  
`set_directive_resource -core RAM_T2P_BRAM array_name`
  - To improve throughput in an array mapped as True dual-Port BRAM

# The source code: my\_defines.h

```
#define MAX_WIDTH          1080
#define MAX_HEIGHT         1920

typedef struct {
    unsigned char R;
    unsigned char G;
    unsigned char B;
} RGB_t;

RGB_t inp_img[MAX_HEIGHT][MAX_WIDTH];
```

# The source code: top\_img\_hist\_equaliz1()

## Design Under Test

```
void top_img_hist_equaliz1( uint11 width, uint11 height,
                           uint25 cdf[GRAY_LEVELS], uint19 hist[GRAY_LEVELS],
                           RGB_t *inp_img,  RGB_t *out_img)
{
    #pragma HLS DATA_PACK variable=out_img
    #pragma HLS INTERFACE ap_fifo port=out_img
    #pragma HLS INTERFACE ap_fifo port=inp_img
    #pragma HLS DATA_PACK variable=inp_img

    /* NOTE: in HW this is a big chunk of memory */
    RGB_t yuv[MAX_WIDTH*MAX_HEIGHT];
    RGB_t yeq[MAX_WIDTH*MAX_HEIGHT];

    // RGB to YUV conversion of current image I(n)
    rgb2yuv(width, height, inp_img, yuv);
    // equalize the current image I(n) and compute its histogram for next image I(n+1)
    img_hist_equaliz1(width, height, hist, cdf, yuv, yeq);
    // YUV to RGB conversion of current image I(n)
    yuv2rgb(width, height, yeq, out_img);
}
```

# The source code: rgb2yuv()

```
void rgb2yuv (uint11 width, uint11 height, RGB_t *in, RGB_t *out) {
    uint11 x, y;
    uint8 R, G, B, Y, U, V;
    const int9 Wrgb[3][3] = { {66,129,25}, {-38,-74,112}, {122,-94,-18} };

    Ly: for (y=0; y<height; y++) {
        Lx: for (x=0; x<width; x++) {

            R = in[y*MAX_WIDTH +x].R;
            G = in[y*MAX_WIDTH +x].G;
            B = in[y*MAX_WIDTH +x].B;

            Y = ( (Wrgb[0][0]*R + Wrgb[0][1]*G + Wrgb[0][2]*B + 128) >> 8) + 16;
            U = ( (Wrgb[1][0]*R + Wrgb[1][1]*G + Wrgb[1][2]*B + 128) >> 8) + 128;
            V = ( (Wrgb[2][0]*R + Wrgb[2][1]*G + Wrgb[2][2]*B + 128) >> 8) + 128;

            out[y*MAX_WIDTH +x].R = Y;
            out[y*MAX_WIDTH +x].G = U;
            out[y*MAX_WIDTH +x].B = V;
        }
    }
}
```



# The source code: yuv2rgb()

```
void yuv2rgb( uint11 width, uint11 height, RGB_t *inp_img, RGB_t *out_img)
{
    int i; uint11 x, y; uint8 R, G, B, Y, U, V; int9 C, D, E;
    const int11 Wyuv[3][3] = { {298, 0, 409}, {298, -100, -208}, {298, 516, 0} };

    Ly: for (y=0; y<height; y++) {
        Lx: for (x=0; x<width; x++) {
            Y=inp_img[y*MAX_WIDTH +x].R; U=inp_img[y*MAX_WIDTH +x].G; V=inp_img[y*MAX_WIDTH +x].B;

            C = Y - 16; D = U - 128; E = V - 128;
            R = CLIP(( Wyuv[0][0] * C + Wyuv[0][2] * E + 128) >> 8);
            G = CLIP(( Wyuv[1][0] * C + Wyuv[1][1] * D + Wyuv[1][2] * E + 128) >> 8);
            B = CLIP(( Wyuv[2][0] * C + Wyuv[2][1] * D + 128) >> 8);

            out_img[y*MAX_WIDTH +x].R = R;
            out_img[y*MAX_WIDTH +x].G = G;
            out_img[y*MAX_WIDTH +x].B = B;

        }
    }
    return;
}
```

# The source code: compute\_cdf()

```
void compute_cdf(uint19 hist[GRAY_LEVELS], uint25 cdf[GRAY_LEVELS])
{
    uint8 i, k;

CDF_L1:    for (k = 0; k < GRAY_LEVELS-1; k++)
    {
        uint25 tmp_cdf=0;
        CDF_L2: for (i = 0; i<=k; i++)
        {
            tmp_cdf = tmp_cdf + hist[i];
        }
        cdf[k] = tmp_cdf;
    }

    return;
}
```

# The source code: img\_hist\_equaliz1

```
void img_hist_equaliz1( uint11 width, uint11 height, uint19 hist[GRAY_LEVELS],
                      uint25 cdf[GRAY_LEVELS], RGB_t *inp_img, RGB_t *out_img)
{
    int i; uint11 x, y; uint8 Y, U, V, cdf_temp; uint32 cdf_mult, gain;
    int numerator = 256*1024*(GRAY_LEVELS-1);
    int denominator = (height*width-1);
    gain =numerator / denominator;
    uint8 old=0, val; uint32 acc=0;
    Ly: for (y=0; y<height; y++) {
        Lx: for (x=0; x<width; x++) {
            Y = inp_img[y*MAX_WIDTH+x].R; U = inp_img[y*MAX_WIDTH+x].G; V = inp_img[y*MAX_WIDTH+x].B;
            //hist[Y]=hist[Y]+1;
            val = Y; if (old==val) acc++; else { hist[old]=acc; acc=hist[val]+1;} old = val;
            cdf_mult = (cdf[Y]-1) * gain ; // histogram equalization
            if (cdf_mult >= (GRAY_LEVELS-1)) cdf_temp = (GRAY_LEVELS-1);
            if (cdf_mult <= 0 )          cdf_temp = 0;
            else                        cdf_temp = (uint8) cdf_mult;
            Y = cdf_temp;
            out_img[y*MAX_WIDTH+x].R = Y; out_img[y*MAX_WIDTH+x].G = U;  out_img[y*MAX_WIDTH+x].B = V;
        } // end of Lx loop
    } // end of Ly loop
    hist[old]=acc;
}
```

# Solution1: synthesis estimation

- `set_directive_loop_tripcount -min 640 -max 1920 func_name/Lx`
- `set_directive_loop_tripcount -min 480 -max 1080 func_name/Ly`

## Summary

Latency		Interval		Type
min	max	min	max	
17513289	120275289	17513290	120275290	none

## Detail

+ Instance

+ Loop

## Utilization Estimates

### Summary

Name	BRAM_18K	DSP48E	FF	LUT
Expression	-	-	-	-
FIFO	-	-	-	-
Instance	-	10	1093	1215
Memory	6144	-	0	0
Multiplexer	-	-	-	144
Register	-	-	9	-
Total	6144	10	1102	1359
Available	280	220	106400	53200
Utilization (%)	2194	4	1	2

**Look at how  
many BRAMs!**

# Solution2: synthesis estimation

## ■ As solution1 plus set\_directive\_dataflow

### [-] Summary

Latency		Interval		Type
min	max	min	max	
17206087	118201687	13210564	91240564	dataflow

### [-] Detail

#### [+] Instance

#### [+] Loop

### Utilization Estimates

#### [-] Summary

Name	BRAM_18K	DSP48E	FF	LUT
Expression	-	-	0	2
FIFO	0	-	20	112
Instance	-	10	1231	1278
Memory	12288	-	0	0
Multiplexer	-	-	-	14
Register	-	-	12	-
Total	12288	10	1263	1406
Available	280	220	106400	53200
Utilization (%)	4388	4	1	2

# solution2 vs. solution1

**Worst performance  
due to double-buffers:**

## Vivado HLS Report Comparison

### Performance Estimates

#### Timing (ns)

Clock		solution1	solution2
default	Target	-	6.25
	Estimated	-	5.47
ap_clk	Target	6.25	-
	Estimated	5.45	-

#### Latency (clock cycles)

		solution1	solution2
Latency	min	17513289	17206087
	max	120275289	118201687
Interval	min	17513290	13210564
	max	120275290	91240564

### Utilization Estimates

	solution1	solution2
BRAM_18K	6144	12288
DSP48E	10	10
FF	1102	1263
LUT	1359	1406

# Solution3: synthesis estimation

- As solution2 plus  
config\_dataflow -  
default\_channel fifo  
-fifo\_depth 2 in the  
script.tcl constraints  
file

**BRAM  
disappeared!**

## Vivado HLS Report Comparison

### All Compared Solutions

[solution3](#): xc7z020clg484-1

### Performance Estimates

#### Timing (ns)

Clock		solution3
default	Target	6.25
	Estimated	5.47

#### Latency (clock cycles)

		solution3
Latency	min	12903363
	max	89166963
Interval	min	12903364
	max	89166964

### Utilization Estimates

	solution3
BRAM_18K	0
DSP48E	10
FF	1160
LUT	1345

# Solution3 vs. solution2

- Small FIFO are much better than double buffers

## Vivado HLS Report Comparison

### Performance Estimates

#### Timing (ns)

Clock		solution3	solution2
default	Target	6.25	6.25
	Estimated	5.47	5.47

#### Latency (clock cycles)

		solution3	solution2
Latency	min	12903363	17206087
	max	89166963	118201687
Interval	min	12903364	13210564
	max	89166964	91240564

### Utilization Estimates

	solution3	solution2
BRAM_18K	0	12288
DSP48E	10	10
FF	1160	1263
LUT	1345	1406



# Solution4: synthesis estimation

- As solution3 plus set\_directive\_pipeline for all Lx loops

## Performance Estimates

### Timing (ns)

#### Summary

Clock	Target	Estimated	Uncertainty
ap_clk	6.25	6.72	0.78

### Latency (clock cycles)

#### Summary

Latency		Interval		Type
min	max	min	max	
614454	4147247	614445	4147245	dataflow

#### Detail

#### Instance

Instance	Module	Latency		Interval		Type
		min	max	min	max	
grp_top_img_hist_equaliz1_rgb2yuv_fu_151	top_img_hist_equaliz1_rgb2yuv	307208	2073608	307208	2073608	none
grp_top_img_hist_equaliz1_img_hist_equaliz1_fu_124	top_img_hist_equaliz1_img_hist_equaliz1	614444	4147244	614444	4147244	none
grp_top_img_hist_equaliz1_yuv2rgb_fu_140	top_img_hist_equaliz1_yuv2rgb	307211	2073611	307211	2073611	none

# solution4 vs. solution3

- Better throughput and latency

## Performance Estimates

### Timing (ns)

Clock		solution4	solution3
default	Target	-	6.25
	Estimated	-	5.47
ap_clk	Target	6.25	-
	Estimated	6.72	-

### Latency (clock cycles)

		solution4	solution3
Latency	min	614454	12903363
	max	4147247	89166963
Interval	min	614445	12903364
	max	4147245	89166964


## Utilization Estimates

	solution4	solution3
BRAM_18K	0	0
DSP48E	12	10
FF	2918	1160
LUT	3160	1345

# Solution5

## ■ As solution4 plus

- set\_directive\_resource -core RAM\_T2P\_BRAM "dpram\_array\_hist" hist
- set\_directive\_dependence -variable hist -type inter -dependent false "dpram\_array\_hist"

Synthesis(solution5) 

[-] **Summary**

Latency		Interval		Type
min	max	min	max	
307255	2073655	307246	2073646	dataflow

[-] **Detail**

[-] **Instance**

Instance	Module	Latency		Interval		Type
		min	max	min	max	
grp_top_img_hist_equaliz1_rgb2yuv_fu_151	top_img_hist_equaliz1_rgb2yuv	307208	2073608	307208	2073608	none
grp_top_img_hist_equaliz1_img_hist_equaliz1_fu_124	top_img_hist_equaliz1_img_hist_equaliz1	307245	2073645	307245	2073645	none
grp_top_img_hist_equaliz1_yuv2rgb_fu_140	top_img_hist_equaliz1_yuv2rgb	307211	2073611	307211	2073611	none

[-] **Loop**

# Solution5 vs solution4

- Better throughput and latency

## Performance Estimates

### Timing (ns)

Clock		solution5	solution4
ap_clk	Target	6.25	6.25
	Estimated	6.72	6.72

### Latency (clock cycles)

		solution5	solution4
Latency	min	307255	614454
	max	2073655	4147247
Interval	min	307246	614445
	max	2073646	4147245

## Utilization Estimates

	solution5	solution4
BRAM_18K	0	0
DSP48E	12	12
FF	2873	2918
LUT	3113	3160

# Summary

## Performance Estimates

### Timing (ns)

Clock		solution1	solution2	solution3	solution4	solution5
default	Target	-	6.25	6.25	-	-
	Estimated	-	5.47	5.47	-	-
ap_clk	Target	6.25	-	-	6.25	6.25
	Estimated	5.45	-	-	6.72	6.72

### Latency (clock cycles)

		solution1	solution2	solution3	solution4	solution5
Latency	min	17513289	17206087	12903363	614454	307255
	max	120275289	118201687	89166963	4147247	2073655
Interval	min	17513290	13210564	12903364	614445	307246
	max	120275290	91240564	89166964	4147245	2073646

## Utilization Estimates

	solution1	solution2	solution3	solution4	solution5
BRAM_18K	6144	12288	0	0	0
DSP48E	10	10	10	12	12
FF	1102	1263	1160	2918	2873
LUT	1359	1406	1345	3160	3113

# A new approach

- Since the 3 functions have the same double FOR loop....
- ...by manually merging all 3 loops into a single one we should save resources and get a better performance

# Exercise on image Histogram equalization

- IN ORDER 1) TO DECREASE RESOURCES AND 2) TO AVOID USING DATA FLOW, WE CAN MANUALLY MERGE (INLINE) THE 3 INNER FUNCTIONS: `rgb2yuv()`, `img_hist_equaliz1()`, `yuv2rgb()` INTO THE TOP LEVEL `top_img_hist_equaliz1`
- YOU NEED ONLY 2 LOOPS: `Ly` and `Lx`
- Steps:
  - 1) Copy `src3` directory into `src4` new directory
  - 2) edit the files into `src4` to manually merge the three functions
  - 3) create a new HLS project: targeting again ZC702 160MHz
  - 4) optimize with PIPELINE and DEPENDENCE as you did for `src3`
  - 5) compares best results of the two different HLS projects

# Synthesis estimation of new approach

Vivado HLS - hls\_hist4\_prj (C:\BagniD\seminar\HLSjul15\5\_img\_hist\_eq\hls\_hist4\_prj) EN English (Ireland)

File Edit Project Solution Window Help

Explorer

- hls\_hist4\_prj
  - Includes
  - Source
  - Test Bench
  - solution1
  - solution2**
    - constraints
      - directives.tcl
      - script.tcl
    - syn
      - report**
        - top\_img\_hist\_equaliz1\_csyntr
      - systemc
      - verilog
      - vhdl

Synthesis(solution2)

**Summary**

Clock	Target	Estimated	Uncertainty
ap_clk	6.25	5.42	0.78

**Latency (clock cycles)**

**Summary**

Latency		Interval		Type
min	max	min	max	Type
307251	2073651	307252	2073652	none

**Detail**

- Instance
- Loop

**Utilization Estimates**

**Summary**

Name	BRAM_18K	DSP48E	FF	LUT
Expression	-	-	0	476
FIFO	-	-	-	-
Instance	-	10	2000	2000
Memory	-	-	-	-
Multiplexer	-	-	-	186
Register	-	-	588	62
<b>Total</b>	<b>0</b>	<b>10</b>	<b>2588</b>	<b>2724</b>
Available	280	220	106400	53200
<b>Utilization (%)</b>	<b>0</b>	<b>4</b>	<b>2</b>	<b>5</b>



# Synthesis performance summary (top\*equaliz1 & top\*equaliz2)

version	FF	LUT	BRAM	DSP48	worst latency	directives
orig. code: top_img_hist_equaliz1						
solution1	1102	1359	6144	10	120275290	set the <a href="#">min</a> and <a href="#">max tripcount</a> for all the loops of all funcs
solution2	1263	1406	12288	10	91240564	as solution 1, plus <a href="#">set_directive_dataflow</a> for top_img_hist_equaliz1()
solution3	1160	1345	0	10	89166964	as solution 2, plus <a href="#">config_dataflow -default_channel fifo -fifo_depth 2</a>
solution4	2918	3160	0	12	4147245	as solution3, plus <a href="#">set_directive_pipeline</a> for all Lx loops of all funcs
solution5	2873	3133	0	12	2073646	as solution4, plus <a href="#">set_directive_dependence</a> for dpram_array_hist()
new code: top_img_hist_equaliz2						
solution	2588	2724	0	10	2073651	all above directives

# Better optimization for CDF(): original

```
#define OT_OPTIMIZED

#ifndef OT_OPTIMIZED
void compute_cdf(uint19 hist[GRAY_LEVELS], uint25 cdf[GRAY_LEVELS])
{
    uint8 i, k;
    CDF_L1: for (k = 0; k < GRAY_LEVELS-1; k++) {
        uint25 tmp_cdf=0;
        CDF_L2: for (i = 0; i<=k; i++) {
            #pragma HLS loop_tripcount min=1 max=256 avg=256
            // cdf for current picture from past picture histogram
            tmp_cdf = tmp_cdf + hist[i];
        }
        cdf[k] = tmp_cdf;
    }
    return;
}
```

# Better optimization for CDF(): optimized

```
#else
void compute_cdf(uint19 hist[GRAY_LEVELS], uint25 cdf[GRAY_LEVELS])
{
    uint8 k;
    uint25 tmp_cdf=0;

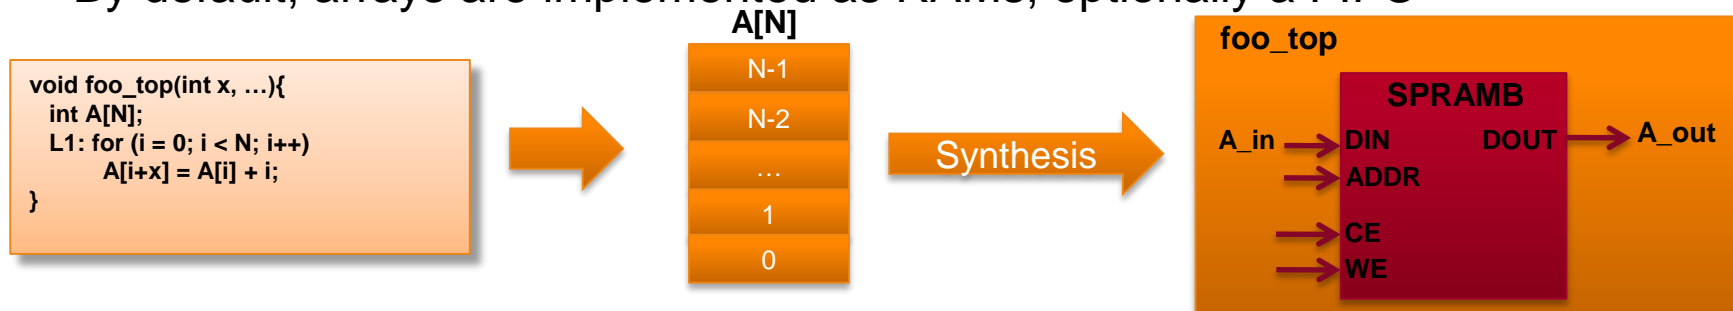
    CDF_L1:for (k = 0; k < GRAY_LEVELS-1; k++)
    {
        #pragma HLS pipeline II=1
        // cdf for current picture from past picture histogram
        tmp_cdf = tmp_cdf + hist[k];
        cdf[k] = tmp_cdf;
    }
    return;
}
#endif
```

# Some directives suitable for arrays

# Review: Arrays in HLS

- An array in C code is implemented by a memory in the RTL

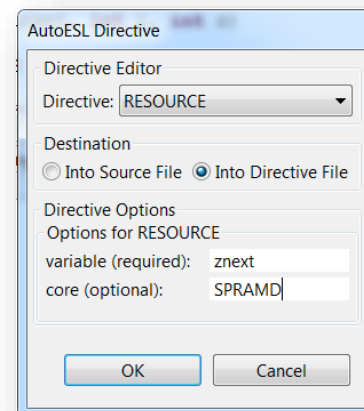
- By default, arrays are implemented as RAMs, optionally a FIFO



- The array can be targeted to any memory resource in the library

- The ports and sequential operation are defined by the library model
  - All RAMs are listed in the AutoESL Library Guide

**Example: Array “znext” is targeted to a single port distributed RAM resource**



# Vertical Array Mapping

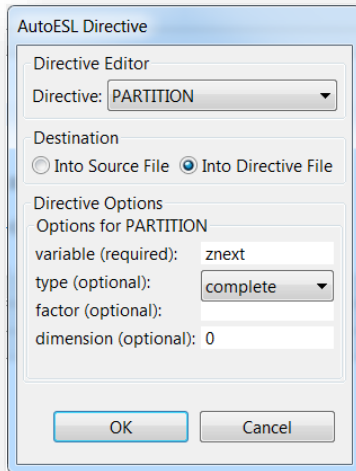
- In vertical mapping, arrays with same length are concatenated by to produce an array with larger bit-widths.
  - i.e. we could move from 2 BRAM units of 1K (depth) x 18bit words each one to only 1 BRAM unit of 1K (depth) x 36bit
- We have seen the effect in the image histogram application:
  - `set_directive_array_map -instance name -mode vertical "func_name" array1_name`
  - `set_directive_array_map -instance name -mode vertical "func_name" array2_name`

# Horizontal Array Mapping

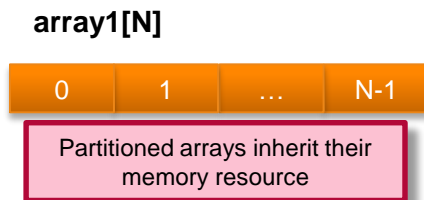
- In horizontal mapping, small arrays are concatenated one after the other to produce a longer array with same bit-widths.
  - `set_directive_array_map -instance name -mode horizontal "func_name" array1_name`
  - `set_directive_array_map -instance name -mode horizontal "func_name" array2_name`
- Although horizontal mapping can result in using less RAM components and hence improve area, it can have an impact on throughput and performance

# Array Partitioning

## ■ Partitioning breaks an array into smaller elements



- Arrays can be split along any dimension
  - If none is specified dimension zero is assumed
  - Dimension zero means all dimensions



Multiple memories allows greater parallel access

block




Divided into blocks:  
N-1/factor elements

cyclic



Divided into blocks:  
1 word at a time  
(like "dealing cards")

complete



Individual elements:  
Break a RAM into registers (no "factor" supported)